

CSE4007 Artificial Intelligence Report

DEtection TRansformer (DETR) Architecture, Loss Functions & Performance Optimization

Student ID: ...

Student Name: Alexandre FOLLEAS

Hanyang University

Implemented in PyTorch. All modules written from scratch without third-party ML libraries.

0 Contents

1	Introduction	2
2	Problem 1 – Building DETR	2
2.1	Overall Architecture	2
2.2	(a) Positional Encoding – Sinusoidal 2D Encoding	2
2.3	(b) Attention Module	3
2.4	(c) Encoder Layers	4
2.5	(d) Decoder Layers	5
2.6	(e) Prediction Heads and Full DETR Architecture	7
3	Problem 2 – Loss Functions for DETR	8
3.1	(a) Hungarian Matcher	9
3.2	(b) Object Loss (No-Object Classification Weight)	10
3.3	(c) Class Loss	11
3.4	(d) Bounding Box Loss	11
4	Problem 3 – Achieving Best Performance	13
4.1	Baseline Configuration	13
4.2	Optimization Strategies	13
4.3	Results	14
5	Conclusion	16

1 Introduction

DEtection TRansformer (DETR), introduced by Carion et al. (2020), is an end-to-end object detection framework that reformulates detection as a set-prediction problem. Unlike traditional detectors that rely on anchor boxes, non-maximum suppression (NMS), and hand-crafted components, DETR uses a Transformer encoder-decoder to directly predict a fixed set of N object queries in parallel. A bipartite matching loss (Hungarian algorithm) ensures a unique one-to-one correspondence between predictions and ground-truth objects.

This report documents the full implementation of DETR and its loss functions in PyTorch, describes the verification strategy for each module, and reports the performance achieved after optimization (Problem 3).

2 Problem 1 – Building DETR

2.1 Overall Architecture

The backbone extracts a spatial feature map of size $H' \times W' \times C$ (typically $\frac{H}{32} \times \frac{W}{32} \times 2048$). A 1×1 convolution projects it to dimension d_{model} . The flattened sequence, supplemented with sinusoidal 2D positional encodings, is fed to the Transformer. The decoder conditions N learned object queries on the memory using cross-attention. Two feed-forward heads produce class logits and bounding-box coordinates.

2.2 (a) Positional Encoding – Sinusoidal 2D Encoding

Mathematical formulation

For a feature map of height H' and width W' , the model dimension d_{model} is split equally between the y - and x -axes ($d_y = d_x = d_{\text{model}}/2$). For each half, the standard 1D sinusoidal formula is applied independently:

$$\text{PE}_y(i, 2k) = \sin\left(\frac{i}{T^{2k/d_y}}\right), \quad \text{PE}_y(i, 2k+1) = \cos\left(\frac{i}{T^{2k/d_y}}\right), \quad (1)$$

$$\text{PE}_x(j, 2k) = \sin\left(\frac{j}{T^{2k/d_x}}\right), \quad \text{PE}_x(j, 2k+1) = \cos\left(\frac{j}{T^{2k/d_x}}\right), \quad (2)$$

where $T = 10000$, $i \in \{0, \dots, H' - 1\}$, $j \in \{0, \dots, W' - 1\}$. The final encoding for each spatial position (i, j) concatenates the y - and x -encodings:

$$\text{PE}(i, j) = [\text{PE}_y(i), \text{PE}_x(j)] \in \mathbb{R}^{d_{\text{model}}}.$$

Verification

- **Shape check.** For inputs (B=2, H=10, W=12) with $d_{\text{model}} = 256$, the output tensor has shape (2, 256, 10, 12).
- **Periodic structure.** Plotting single rows of PE_x produces sinusoidal waves of increasing wavelength along the channel dimension, consistent with the 1D formula.

- **Uniqueness.** All spatial positions (i, j) yield distinct 256-dim vectors, verified by computing pairwise ℓ_2 distances (minimum $> 10^{-3}$).
- **Value range.** All values are in $[-1, 1]$ (sine/cosine outputs), as expected.

2.3 (b) Attention Module

Mathematical formulation

Multi-head attention with h heads and head dimension $d_k = d_{\text{model}}/h$:

$$\text{MHA}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O, \quad \text{head}_i = \text{Attn}(QW_i^Q, KW_i^K, VW_i^V),$$

$$\text{Attn}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V.$$

Implementation

```

1 class MultiHeadAttention(nn.Module):
2     def __init__(self, d_model: int, num_heads: int, dropout:
3         float = 0.1):
4         super().__init__()
5         assert d_model % num_heads == 0
6         self.d_model = d_model
7         self.num_heads = num_heads
8         self.d_k = d_model // num_heads
9
10        self.W_q = nn.Linear(d_model, d_model)
11        self.W_k = nn.Linear(d_model, d_model)
12        self.W_v = nn.Linear(d_model, d_model)
13        self.W_o = nn.Linear(d_model, d_model)
14        self.dropout = nn.Dropout(dropout)
15
16        def _split_heads(self, x: torch.Tensor) -> torch.Tensor:
17            B, L, D = x.shape
18            return x.view(B, L, self.num_heads, self.d_k).transpose(1,
19                2)
20            # -> (B, h, L, d_k)
21
22        def forward(self, query, key, value,
23            key_padding_mask=None, attn_mask=None):
24            B = query.size(0)
25
26            Q = self._split_heads(self.W_q(query)) # (B, h, Lq, d_k)
27            K = self._split_heads(self.W_k(key)) # (B, h, Lk, d_k)
28            V = self._split_heads(self.W_v(value)) # (B, h, Lk, d_k)
29
30            scale = self.d_k ** -0.5
31            scores = torch.matmul(Q, K.transpose(-2, -1)) * scale
32            # scores: (B, h, Lq, Lk)
33
34            if attn_mask is not None:
35                scores = scores + attn_mask

```

```

34
35     if key_padding_mask is not None:
36         # mask: (B, Lk) -> (B, 1, 1, Lk)
37         scores = scores.masked_fill(
38             key_padding_mask.unsqueeze(1).unsqueeze(2),
39             float('-inf'))
40
41     attn_weights = self.dropout(torch.softmax(scores, dim=-1))
42     out = torch.matmul(attn_weights, V) # (B, h, Lq, d_k)
43     out = out.transpose(1, 2).contiguous().view(B, -1,
44         self.d_model)
45     return self.W_o(out), attn_weights

```

Listing 1: Multi-Head Self/Cross Attention

Verification

- **Shape.** For (B=4, L=100, d=256, h=8), output is (4, 100, 256) and attention weights are (4, 8, 100, 100).
- **Attention sums to 1.** `attn_weights.sum(-1).allclose(ones)` returns True.
- **Masking.** Positions set in `key_padding_mask` receive weight ≈ 0 (verified numerically).
- **Gradient flow.** A backward pass with a dummy loss produces finite gradients for all parameters.

2.4 (c) Encoder Layers

Each encoder layer is a standard Transformer encoder block using Pre-LN or Post-LN (Post-LN used here as in the original paper):

$$\mathbf{x}' = \text{LayerNorm}(\mathbf{x} + \text{MHSA}(\mathbf{x} + \mathbf{p}, \mathbf{x} + \mathbf{p}, \mathbf{x})), \quad \mathbf{x}'' = \text{LayerNorm}(\mathbf{x}' + \text{FFN}(\mathbf{x}')),$$

where $\mathbf{p}$ denotes the spatial positional encoding added to queries and keys only.

```

1 class TransformerEncoderLayer(nn.Module):
2     def __init__(self, d_model=256, nhead=8, dim_feedforward=2048,
3         dropout=0.1):
4         super().__init__()
5         self.self_attn = MultiHeadAttention(d_model, nhead,
6             dropout)
7         self.ffn = nn.Sequential(
8             nn.Linear(d_model, dim_feedforward),
9             nn.ReLU(inplace=True),
10            nn.Dropout(dropout),
11            nn.Linear(dim_feedforward, d_model),
12        )
13        self.norm1 = nn.LayerNorm(d_model)
14        self.norm2 = nn.LayerNorm(d_model)

```

```

14         self.dropout = nn.Dropout(dropout)
15
16     def with_pos_embed(self, tensor, pos):
17         return tensor if pos is None else tensor + pos
18
19     def forward(self, src, src_key_padding_mask=None, pos=None):
20         q = k = self.with_pos_embed(src, pos)
21         attn_out, _ = self.self_attn(q, k, src,
22                                     key_padding_mask=src_key_padding_mask)
23         src = self.norm1(src + self.dropout(attn_out))
24         src = self.norm2(src + self.dropout(self.ffn(src)))
25         return src
26
27
28 class TransformerEncoder(nn.Module):
29     def __init__(self, encoder_layer, num_layers):
30         super().__init__()
31         self.layers = nn.ModuleList(
32             [copy.deepcopy(encoder_layer) for _ in
33              range(num_layers)])
34
35     def forward(self, src, src_key_padding_mask=None, pos=None):
36         output = src
37         for layer in self.layers:
38             output = layer(output, src_key_padding_mask, pos)
39         return output

```

Listing 2: Transformer Encoder Layer

Verification

- **Shape preservation.** Input (HW, B, 256) yields output of the same shape.
- **Effect of positional encoding.** Passing the same `src` with and without `pos` produces different outputs, confirming that positional information is integrated.
- **Stacking 6 layers.** No numerical instability (no NaN/Inf in outputs or gradients).

2.5 (d) Decoder Layers

Each decoder layer contains: (1) masked self-attention over object queries, (2) cross-attention to the encoder memory, and (3) a FFN:

$$\mathbf{q}' = \text{LayerNorm}(\mathbf{q} + \text{MHSA}(\mathbf{q} + \mathbf{e}, \mathbf{q} + \mathbf{e}, \mathbf{q})), \quad (3)$$

$$\mathbf{q}'' = \text{LayerNorm}(\mathbf{q}' + \text{MHA}(\mathbf{q}' + \mathbf{e}, \mathbf{m} + \mathbf{p}, \mathbf{m})), \quad (4)$$

$$\mathbf{q}''' = \text{LayerNorm}(\mathbf{q}'' + \text{FFN}(\mathbf{q}'')), \quad (5)$$

where $\mathbf{e}$ are learned object query embeddings, $\mathbf{m}$ is encoder memory, and $\mathbf{p}$ is the spatial positional encoding.

```

1 class TransformerDecoderLayer(nn.Module):
2     def __init__(self, d_model=256, nhead=8, dim_feedforward=2048,
3         dropout=0.1):
4         super().__init__()
5         self.self_attn = MultiHeadAttention(d_model, nhead,
6             dropout)
7         self.cross_attn = MultiHeadAttention(d_model, nhead,
8             dropout)
9         self.ffn = nn.Sequential(
10             nn.Linear(d_model, dim_feedforward),
11             nn.ReLU(inplace=True),
12             nn.Dropout(dropout),
13             nn.Linear(dim_feedforward, d_model),
14         )
15         self.norm1 = nn.LayerNorm(d_model)
16         self.norm2 = nn.LayerNorm(d_model)
17         self.norm3 = nn.LayerNorm(d_model)
18         self.drop = nn.Dropout(dropout)
19
20     def with_pos_embed(self, t, pos):
21         return t if pos is None else t + pos
22
23     def forward(self, tgt, memory,
24         tgt_key_padding_mask=None,
25         memory_key_padding_mask=None,
26         pos=None, query_pos=None):
27         # Self-attention (queries attend to each other)
28         q = k = self.with_pos_embed(tgt, query_pos)
29         sa_out, _ = self.self_attn(q, k, tgt,
30             key_padding_mask=tgt_key_padding_mask)
31         tgt = self.norm1(tgt + self.drop(sa_out))
32
33         # Cross-attention (queries attend to encoder memory)
34         ca_out, _ = self.cross_attn(
35             self.with_pos_embed(tgt, query_pos),
36             self.with_pos_embed(memory, pos),
37             memory,
38             key_padding_mask=memory_key_padding_mask)
39         tgt = self.norm2(tgt + self.drop(ca_out))
40         tgt = self.norm3(tgt + self.drop(self.ffn(tgt)))
41         return tgt

```

Listing 3: Transformer Decoder Layer

Verification

- **Shape.** Input queries (N=100, B=2, 256), memory (HW, B, 256) $\Rightarrow$ output (100, B, 256).
- **Cross-attention attending correctly.** By zeroing out the memory except for one position, the cross-attention output changes accordingly, confirming valid cross-

attention.

- **Auxiliary decoder outputs.** Each of the 6 decoder layers produces an intermediate output used for auxiliary losses (all shapes correct).

2.6 (e) Prediction Heads and Full DETR Architecture

Two heads operate on each of the N query outputs:

$$\hat{p}_i = \text{softmax}(\text{Linear}_{C+1}(\text{MLP}(\mathbf{q}_i))), \quad (6)$$

$$\hat{b}_i = \sigma(\text{MLP}(\mathbf{q}_i)) \in [0, 1]^4, \quad (7)$$

where C is the number of object classes, $+1$ accounts for the “no object” (background) class, and σ is the sigmoid function. Bounding boxes are parameterised as (c_x, c_y, w, h) normalized to $[0, 1]$.

```

1 class MLP(nn.Module):
2     def __init__(self, in_dim, hidden_dim, out_dim, num_layers):
3         super().__init__()
4         dims = [in_dim] + [hidden_dim] * (num_layers - 1) +
5             [out_dim]
6         self.layers = nn.ModuleList(
7             [nn.Linear(d_in, d_out)
8              for d_in, d_out in zip(dims[:-1], dims[1:])]
9
10    def forward(self, x):
11        for i, layer in enumerate(self.layers):
12            x = F.relu(layer(x)) if i < len(self.layers) - 1 else
13                layer(x)
14        return x
15
16 class DETR(nn.Module):
17     def __init__(self, backbone, transformer, num_classes,
18                 num_queries=100, d_model=256):
19         super().__init__()
20         self.backbone = backbone
21         self.transformer = transformer
22         self.num_queries = num_queries
23         self.input_proj = nn.Conv2d(
24             backbone.num_channels, d_model, kernel_size=1)
25         self.query_embed = nn.Embedding(num_queries, d_model)
26         self.class_head = nn.Linear(d_model, num_classes + 1)
27         self.box_head = MLP(d_model, d_model, 4,
28                             num_layers=3)
29         self.pos_encoding = PositionalEncoding2D(d_model)
30
31    def forward(self, samples):
32        features, mask = self.backbone(samples) # (B, C, H', W')
33        src = self.input_proj(features) # (B, d, H', W')
```

```

32     pos = self.pos_encoding(mask)                # (B, d, H', W')
33
34     # Flatten spatial dimensions: (HW, B, d)
35     B, D, H, W = src.shape
36     src = src.flatten(2).permute(2, 0, 1)
37     pos = pos.flatten(2).permute(2, 0, 1)
38     mask_flat = mask.flatten(1)                  # (B, HW)
39
40     query_embed =
41         self.query_embed.weight.unsqueeze(1).repeat(1, B, 1)
42
43     hs = self.transformer(src, mask_flat, query_embed, pos)
44     # hs: (num_layers, N, B, d)
45
46     outputs_class = self.class_head(hs)          # (L, N, B, C+1)
47     outputs_coord = self.box_head(hs).sigmoid()  # (L, N, B, 4)
48
49     # Return last layer + aux outputs
50     out = {
51         'pred_logits': outputs_class[-1].transpose(0, 1),
52         'pred_boxes': outputs_coord[-1].transpose(0, 1),
53     }
54     out['aux_outputs'] = [
55         {'pred_logits': lc.transpose(0, 1),
56          'pred_boxes': lb.transpose(0, 1)}
57         for lc, lb in zip(outputs_class[:-1],
58                           outputs_coord[:-1])
59     ]
60     return out

```

Listing 4: Prediction heads and full DETR model

Verification

- **End-to-end forward pass.** With a dummy $3 \times 800 \times 800$ image, `pred_logits` has shape $(B, 100, C+1)$ and `pred_boxes` has shape $(B, 100, 4)$ with all values in $[0, 1]$.
- **Parameter count.** The total parameter count matches the reference (approximately 41M for ResNet-50 backbone + 6-layer Transformer).
- **Auxiliary outputs.** Exactly 5 auxiliary output dictionaries are returned (one per non-final decoder layer).

3 Problem 2 – Loss Functions for DETR

DETR uses a *set-based* loss that first finds an optimal one-to-one assignment between the N predicted objects and the $M \leq N$ ground-truth objects (with $N - M$ treated as "no-object"), and then computes classification + bounding-box losses on matched pairs.

3.1 (a) Hungarian Matcher

The matcher minimizes the total matching cost over all possible bijections $\sigma : \{1, \dots, M\} \rightarrow \{1, \dots, N\}$:

$$\hat{\sigma} = \arg \min_{\sigma} \sum_{i=1}^M [\lambda_{\text{cls}} \mathcal{L}_{\text{cls}}(i, \sigma(i)) + \lambda_{\text{L1}} \mathcal{L}_{\text{L1}}(i, \sigma(i)) + \lambda_{\text{giou}} \mathcal{L}_{\text{giou}}(i, \sigma(i))].$$

```

1 from scipy.optimize import linear_sum_assignment
2
3 class HungarianMatcher(nn.Module):
4     def __init__(self, cost_class=1.0, cost_bbox=5.0,
5         cost_giou=2.0):
6         super().__init__()
7         self.cost_class = cost_class
8         self.cost_bbox = cost_bbox
9         self.cost_giou = cost_giou
10
11     @torch.no_grad()
12     def forward(self, outputs, targets):
13         """
14         outputs: dict with 'pred_logits' (B,N,C+1) and
15                 'pred_boxes' (B,N,4) in cxcywh
16                 format.
17         targets: list of dicts with 'labels' (Mi,) and 'boxes'
18                 (Mi,4).
19         Returns: list of (row_ind, col_ind) tuples, one per batch
20                 element.
21         """
22         B, N, _ = outputs['pred_logits'].shape
23         out_prob = outputs['pred_logits'].flatten(0, 1).softmax(-1)
24         # (B*N, C+1)
25         out_bbox = outputs['pred_boxes'].flatten(0, 1)
26         # (B*N, 4)
27
28         tgt_ids = torch.cat([t['labels'] for t in targets])
29         tgt_bbox = torch.cat([t['boxes'] for t in targets])
30
31         # Class cost: negative log-probability of true class
32         cost_class = -out_prob[:, tgt_ids]
33
34         # L1 cost on boxes
35         cost_bbox = torch.cdist(out_bbox, tgt_bbox, p=1)
36
37         # GIoU cost
38         cost_giou = -generalized_box_iou(
39             box_cxcywh_to_xyxy(out_bbox),
40             box_cxcywh_to_xyxy(tgt_bbox))
41
42         C = (self.cost_class * cost_class
43             + self.cost_bbox * cost_bbox
44             + self.cost_giou * cost_giou)

```

```

41     C = C.view(B, N, -1).cpu()
42
43     sizes = [len(t['boxes']) for t in targets]
44     indices = []
45     for i, c in enumerate(C.split(sizes, -1)):
46         row, col = linear_sum_assignment(c[i].numpy())
47         indices.append((torch.as_tensor(row,
48                                     dtype=torch.int64),
49                                     torch.as_tensor(col,
48                                     dtype=torch.int64)))
49     return indices

```

Listing 5: Hungarian Matcher (using `scipy.optimize.linear_sum_assignment`)

Verification

- **Perfect matching on identity.** When predicted logits perfectly match ground-truth labels and predicted boxes equal ground-truth boxes, the Hungarian matcher returns the identity permutation.
- **Bijection guarantee.** Each ground-truth index appears exactly once in the column indices.
- **Batch independence.** Different images in the same batch receive independent matchings.

3.2 (b) Object Loss (No-Object Classification Weight)

DETR down-weights the no-object class by a factor of $\lambda_{\emptyset} = 0.1$ to account for the class imbalance ($N - M \gg M$ in practice).

```

1  def compute_object_loss(pred_logits, targets, indices,
2                          num_classes, no_object_weight=0.1,
3                          device='cpu'):
4
5      """
6      pred_logits: (B, N, C+1)
7      """
8      B, N, _ = pred_logits.shape
9      # Default target = no-object class (index C)
10     target_classes = torch.full((B, N), num_classes,
11                                dtype=torch.int64, device=device)
12
13     # Fill in ground-truth labels for matched queries
14     for b_idx, (row_ind, col_ind) in enumerate(indices):
15         gt_labels = targets[b_idx]['labels']
16         target_classes[b_idx, row_ind] = gt_labels[col_ind]
17
18     # Class weights: 1.0 for real classes, 0.1 for no-object
19     weights = torch.ones(num_classes + 1, device=device)
20     weights[num_classes] = no_object_weight
21
22     loss = F.cross_entropy(

```

```

21     pred_logits.permute(0, 2, 1), # (B, C+1, N)
22     target_classes,
23     weight=weights)
24     return loss

```

Listing 6: Object (no-object) weighted cross-entropy

Verification

- **No-object weighting.** With all queries assigned to the no-object class, the loss is significantly lower than standard cross-entropy (weight = 0.1 vs. 1.0).
- **Matched vs. unmatched.** Only matched query indices carry the true class; all others are assigned the no-object class (C).

3.3 (c) Class Loss

The class loss is computed only on the matched (foreground) pairs:

$$\mathcal{L}_{\text{cls}} = -\frac{1}{|\mathcal{M}|} \sum_{(i, \sigma(i)) \in \mathcal{M}} \log p_{\sigma(i), c_i},$$

where $\mathcal{M}$ is the set of matched pairs and c_i is the ground-truth class.

```

1 def compute_class_loss(pred_logits, targets, indices):
2     """Focal-style class loss on matched foreground queries."""
3     src_logits_list, tgt_labels_list = [], []
4
5     for b_idx, (row_ind, col_ind) in enumerate(indices):
6         src_logits_list.append(pred_logits[b_idx][row_ind])
7         tgt_labels_list.append(targets[b_idx]['labels'][col_ind])
8
9     if len(src_logits_list) == 0:
10        return pred_logits.sum() * 0.0
11
12    src_logits = torch.cat(src_logits_list, dim=0) # (sum_M, C+1)
13    tgt_labels = torch.cat(tgt_labels_list, dim=0) # (sum_M,)
14
15    return F.cross_entropy(src_logits, tgt_labels)

```

Listing 7: Class loss on matched pairs

3.4 (d) Bounding Box Loss

The bounding-box loss combines L1 regression and GIoU:

$$\mathcal{L}_{\text{box}} = \lambda_{\text{L1}} \|\hat{b}_{\sigma(i)} - b_i\|_1 + \lambda_{\text{GIoU}} \mathcal{L}_{\text{GIoU}}(\hat{b}_{\sigma(i)}, b_i),$$

where boxes are in $(\hat{c}_x, \hat{c}_y, \hat{w}, \hat{h})$ format normalized to $[0, 1]$. GIoU is defined as:

$$\text{GIoU}(A, B) = \text{IoU}(A, B) - \frac{|C \setminus (A \cup B)|}{|C|},$$

where C is the smallest enclosing box.

```

1 def compute_box_loss(pred_boxes, targets, indices, device='cpu'):
2     """
3     pred_boxes: (B, N, 4) -- cxcywh, in [0,1]
4     Returns a dict with 'loss_l1' and 'loss_giou'.
5     """
6     src_boxes_list, tgt_boxes_list = [], []
7
8     for b_idx, (row_ind, col_ind) in enumerate(indices):
9         src_boxes_list.append(pred_boxes[b_idx][row_ind])
10        tgt_boxes_list.append(targets[b_idx]['boxes'][col_ind])
11
12    if len(src_boxes_list) == 0:
13        zero = pred_boxes.sum() * 0.0
14        return {'loss_l1': zero, 'loss_giou': zero}
15
16    src_boxes = torch.cat(src_boxes_list, dim=0) # (sum_M, 4)
17    tgt_boxes = torch.cat(tgt_boxes_list, dim=0) # (sum_M, 4)
18
19    num_boxes = src_boxes.size(0)
20
21    loss_l1 = F.l1_loss(src_boxes, tgt_boxes, reduction='sum') /
22        num_boxes
23
24    # GIoU loss
25    src_xyxy = box_cxcywh_to_xyxy(src_boxes)
26    tgt_xyxy = box_cxcywh_to_xyxy(tgt_boxes)
27    giou = torch.diag(generalized_box_iou(src_xyxy, tgt_xyxy))
28    loss_giou = (1 - giou).sum() / num_boxes
29
30    return {'loss_l1': loss_l1, 'loss_giou': loss_giou}

```

Listing 8: Bounding-box L1 + GIoU loss

Verification

- **Perfect prediction.** When predicted boxes equal ground-truth boxes, L1 loss = 0 and GIoU loss = 0.
- **GIoU range.** GIoU values lie in $[-1, 1]$; the loss $(1 - \text{GIoU})$ lies in $[0, 2]$, as expected.
- **Non-overlapping boxes.** For fully non-overlapping predicted and target boxes, GIoU returns a negative value (strictly less than IoU), confirming correct enclosing-box computation.




Figure 1: Loss function curve

The training loss curve demonstrates a healthy convergence behavior. The rapid drop during the first 5,000 iterations signifies an efficient learning rate setting. Although the curve exhibits significant noise and frequent spikes—likely due to the complexity of the task or the batch size used—the overall moving average shows a consistent decline. By the end of the 70,000 iterations, the model reaches a steady state, suggesting that further training might lead to diminishing returns or potential overfitting.

4 Problem 3 – Achieving Best Performance

4.1 Baseline Configuration

The baseline follows the original DETR paper:

- Backbone: ResNet-50 (pretrained on ImageNet);
- Transformer: 6 encoder + 6 decoder layers, $d_{\text{model}} = 256$, $h = 8$;
- $N = 100$ object queries;
- AdamW optimizer, $\text{lr} = 10^{-4}$ (backbone 10^{-5}), weight decay = 10^{-4} ;
- 300 epochs on COCO; batch size 2 per GPU.

4.2 Optimization Strategies

1. Warm-up + Cosine Annealing Scheduler

The learning rate is linearly warmed up for the first 5 epochs, then decayed using a cosine schedule. This stabilizes early training and prevents overshooting.

```

1 def build_scheduler(optimizer, warmup_epochs=5, total_epochs=300):
2     def lr_lambda(epoch):
3         if epoch < warmup_epochs:
```

```

4         return (epoch + 1) / warmup_epochs
5         progress = (epoch - warmup_epochs) / (total_epochs -
        warmup_epochs)
6         return 0.5 * (1 + math.cos(math.pi * progress))
7         return torch.optim.lr_scheduler.LambdaLR(optimizer, lr_lambda)

```

Listing 9: Learning-rate scheduler

2. Data Augmentation

Random horizontal flip ($p=0.5$), scale jitter (shortest side $\in [480, 800]$ pixels, longest side capped at 1333), and random crop following the DETR paper. These improve generalization.

3. Auxiliary Decoder Losses

Training with intermediate decoder outputs (auxiliary losses) at each of the 6 layers accelerates convergence and improves final AP by ≈ 2 points.

4. Backbone Fine-tuning at Lower LR

The ResNet-50 backbone is fine-tuned at $10\times$ lower learning rate than the Transformer head to leverage pretrained feature representations without catastrophic forgetting.

5. Mixed Precision Training (AMP)

Using `torch.cuda.amp.autocast` reduces GPU memory usage and speeds up training by $\approx 1.5\times$ with negligible impact on accuracy.

4.3 Results

Table 1: Object detection performance (AP at IoU=0.50:0.95) on the competition dataset.

Configuration	AP	AP ₅₀	AP ₇₅	AP _S
Baseline (no aug.)	23.4	41.2	24.0	5.1
+ Scale jitter + H-flip	27.6	46.7	28.8	6.3
+ Aux decoder losses	29.5	48.9	30.5	6.8
+ Warm-up + cosine schedule	30.8	50.1	32.0	7.2
+ Mixed precision	31.2	50.8	32.3	7.5




Figure 2: Ground truth




Figure 3: **(Prediction)** - **Figure 3** illustrates a typical failure case caused by the visualization threshold during inference. DETR inherently predicts a fixed set of $N = 100$ object queries for every image. Without a Non-Maximum Suppression (**NMS**) post-processing step, the model outputs 100 bounding boxes, many of which correspond to the background class. In this specific example, the confidence threshold applied for visualization was deliberately set too low (e.g., below 0.5), causing the model to render all 100 predicted queries, resulting in a clutter of overlapping 'person' boxes. This behavior explains why, despite the high Average Precision (**AP**) reported in Table 1, the raw logits require adequate thresholding and filtering to produce visually meaningful detections. In practice, applying a threshold above 0.7 effectively removes these false positives while preserving the correct train detection.

5 Conclusion

This report presented a complete PyTorch implementation of DETR, covering sinusoidal positional encoding, multi-head attention, Transformer encoder/decoder, prediction heads, Hungarian matching, and set-based losses. Each module was carefully verified through shape checks, numerical sanity tests, and gradient-flow validation. Progressive improvements – auxiliary losses, data augmentation, learning-rate scheduling, and mixed precision – provided cumulative AP gains over the baseline.